

ASSIGNMENT:- 1

AIM:- Linear Regression

PROBLEM STATEMENT:- Real estate agents want help to predict the house price for regions in the USA. He gave you the dataset to work on and you decided to use the Linear Regression Model. Create a model that will help him to estimate what the house would sell for.

URL for the dataset: https://github.com/huzaihsayed/Linear-Regression-Model-for-House-PricePrediction/blob/master/USA_Housing.csv

THEORY:-

Regression Analysis

- Regression analysis is one of the core techniques in supervised machine learning used to model relationships between a dependent variable (target) and one or more independent variables (features).
- The primary goal is to predict the target variable using an optimized function derived from the independent variables.

Linear Regression


Linear Regression is a supervised learning algorithm used to model the linear relationship between independent variables (features) and a dependent variable (target). It assumes the relationship can be expressed as:

The diagram shows the linear regression equation $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \varepsilon$ with arrows pointing to each term from descriptive labels. An arrow from 'Dependent Variable (Response Variable)' points to Y . An arrow from 'Independent Variables (Predictors)' points to the X terms. An arrow from 'Y intercept' points to β_0 . An arrow from 'Slope Coefficient' points to β_1 and β_2 . An arrow from 'Error Term' points to ε .

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \varepsilon$$

Labels and their corresponding parts in the equation:

- Dependent Variable (Response Variable) → Y
- Independent Variables (Predictors) → X_1, X_2, \dots
- Y intercept → β_0
- Slope Coefficient → β_1, β_2
- Error Term → ε



Application to Housing Data

The dataset includes features like Avg. Area Income, Avg. Area House Age, and Area Population. Linear Regression will estimate how these features collectively influence house prices.

ADVANTAGES:-

- Simple to implement and interpret.
- Computationally efficient.
- Provides insights into feature importance through coefficients.

LIMITATIONS:-

- Assumes linear relationships; fails to capture non-linear patterns.
- Sensitive to outliers and multicollinearity.

ASSUMPTIONS:-

- Linearity: Relationship between features and target is linear.
- Independence: Observations are independent of each other.
- Homoscedasticity: Residuals have constant variance across predictions.
- Normality: Residuals are normally distributed.
- No Multicollinearity: Features are not highly correlated with each other.

EVALUATION METRICS:-

Mean Absolute Error (MAE):

Average absolute difference between predictions and actual values.

Interpretability: Represents the average absolute error in the same units as the target variable (e.g., dollars for house prices).

Robustness: Less sensitive to outliers compared to MSE, making it useful when extreme errors should not dominate the evaluation.

Mean Squared Error (MSE):

Average squared difference penalizes larger errors.

Penalizes Larger Errors: Squaring the errors gives higher weight to significant deviations (e.g., a \$200K error contributes more than a \$50K error).

Downside: Sensitive to outliers—large errors can disproportionately increase MSE.

Root Mean Squared Error (RMSE):

Square root of MSE, in the same units as the target.

Same Units as Target: Unlike MSE (which is in squared units), RMSE is in the original units (e.g., dollars), making it easier to interpret.

Industry Standard: Commonly used in real estate and finance for evaluating prediction models.


OUTPUT ANALYSIS:-

Metrics:

MAE: 82,288

MSE: 10,460,958,907

RMSE: 102,278



In the above histogram plot, we see data is in bell shape (Normally Distributed), which indicates that the model has accurate predictions.

CONCLUSION:-

The project used Linear Regression, a foundational statistical and machine learning algorithm, to model and predict house prices based on various features. The assumption that there exists a linear relationship between these features and the house price seems valid, as seen from the normally distributed residuals and the model performance metrics.

The model achieved the following results that suggest:

On average, the model's predictions are off by around \$82,288, which is reasonable depending on the scale of home prices in the dataset.

The residuals (errors between predicted and actual prices) are normally distributed, as indicated by the bell-shaped histogram, which means the model satisfies one of the core linear regression assumptions and performs reliably on the given data.

The model can be used by real estate agents, property investors, or buyers/sellers to estimate the fair market value of a property.

ASSIGNMENT:- 2

AIM:- Multiclass Classification

PROBLEM STATEMENT:- Build a Multiclass classifier using the CNN model. Use MNIST or any other suitable dataset. a. Perform Data Pre-processing b. Define Model and perform training c. Evaluate Results using confusion matrix.

THEORY:-

Convolutional Neural Network (CNN)

Convolutional Neural Networks (CNNs) are a class of deep learning models specifically designed to process and analyze data with a grid-like topology, such as images. They are widely used for tasks like image classification, object detection, and facial recognition. CNNs are especially effective for multiclass classification problems where the goal is to assign each input image to one of several predefined categories.

A CNN typically consists of multiple layers, including convolutional layers, pooling layers, and fully connected layers. Convolutional layers extract important features from input images using learnable filters, while pooling layers reduce the spatial dimensions, leading to lower computational costs and better generalization. The fully connected layers at the end perform the final classification based on the extracted features.

Description of the MNIST Dataset

The MNIST (Modified National Institute of Standards and Technology) dataset is a benchmark dataset commonly used for training and evaluating image classification models. It consists of 70,000 grayscale images of handwritten digits (0 through 9), each of size 28x28 pixels. The dataset is divided into:

- Training set: 60,000 images
- Test set: 10,000 images

Each image is labeled with the digit it represents, making it ideal for a 10-class classification problem. The simplicity and standardization of the MNIST dataset make it a suitable choice for evaluating the performance of CNNs on multiclass classification tasks.

a. Data Pre-processing

Data pre-processing is a crucial step in building an effective CNN model. For the MNIST dataset, this involves the following tasks:

- Normalization: Pixel values ranging from 0 to 255 are scaled to a range of 0 to 1 to improve model convergence.
- Reshaping: The input images are reshaped to match the CNN input requirements (e.g., (28, 28, 1) for grayscale images).
- One-hot Encoding: The class labels (0–9) are converted to one-hot encoded vectors to be compatible with the softmax output layer.

b. Defining the CNN Model and Performing Training

The CNN model is defined using layers that extract spatial features and perform classification. A basic CNN architecture for the MNIST dataset may include:

- Convolutional Layers: To extract local patterns using filters.
- ReLU Activation: To introduce non-linearity.
- Max Pooling Layers: To reduce dimensionality while retaining essential features.
- Flatten Layer: To convert the 2D feature maps into a 1D vector.
- Fully Connected (Dense) Layers: To perform classification based on extracted features.


The model is compiled using a suitable optimizer (e.g., Adam) and loss function (e.g., categorical crossentropy). Training is done over multiple epochs with batch processing to minimize the loss and improve accuracy.

Evaluating results of the Model

The effectiveness of the CNN model is measured using the following evaluation metrics.

1. Confusion Matrix

For the MNIST dataset, the confusion matrix will be a 10x10 matrix, since there are ten digits.



2. Accuracy

```
313/313 [=====] - 3s 8ms/step - loss: 0.0286 - accuracy: 0.9894  
Test Accuracy: 98.94%
```

3. Precision, Recall, and F1-Score

- **Precision:** Measures the proportion of true positive predictions out of all positive predictions made.

$$\text{Precision}_{\text{Class A}} = \frac{TP_{\text{Class A}}}{TP_{\text{Class A}} + FP_{\text{Class A}}}$$



- **Recall:** Measures the proportion of true positives out of all actual positives.

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$



- **F1-Score:** The harmonic mean of precision and recall.

$$F1 = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

Classification Report:

Classification Report:				
	precision	recall	f1-score	support
0	0.99	1.00	0.99	980
1	0.99	1.00	0.99	1135
2	0.99	0.99	0.99	1032
3	0.99	0.99	0.99	1010
4	0.99	0.99	0.99	982
5	0.99	0.98	0.99	892
6	0.99	0.99	0.99	958
7	0.98	0.99	0.99	1028
8	0.98	0.99	0.99	974
9	0.99	0.98	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

Advantages of CNN

- **Automatic Feature Extraction:** CNNs automatically learn relevant features from raw image data during training, eliminating the need for manual feature engineering.
- **Spatial Invariance:** Through the use of convolution and pooling layers, CNNs are effective at recognizing patterns regardless of their position in the input image.
- **High Accuracy on Image Data:** CNNs consistently achieve high accuracy for image classification tasks, especially with large and complex datasets such as MNIST, CIFAR, or ImageNet.

Limitations of CNN

- **Computationally Intensive:** Training CNNs requires significant computational power, especially with deep architectures and large datasets, often requiring GPUs.

- **Large Amount of Labeled Data Required:** CNNs typically require a large amount of labeled data to perform well, which may not always be available for every application.
- **Lack of Interpretability:** CNNs function as black-box models, making it difficult to interpret or explain the reasoning behind a specific prediction.

CONCLUSION:-

This assignment involved implementing a Convolutional Neural Network (CNN) to classify handwritten digits from the MNIST dataset. The CNN architecture successfully extracted meaningful features from 28x28 grayscale images and learned to distinguish digits from 0 to 9 with high accuracy. This validates the strength of CNNs in visual recognition tasks, especially where spatial hierarchies and local patterns are critical.

The following preprocessing steps significantly improved the model's performance: Normalization, Reshaping and One-hot Encoding

The model's layered design provided insights into how CNNs operate:

Convolutional Layers captured spatial patterns (like curves, edges) in digits.

ReLU Activation introduced non-linearity, enabling complex representations.

Pooling Layers reduced dimensionality, preventing overfitting and increasing computational efficiency.

Fully Connected Layers + Softmax translated abstracted features into class probabilities.

Accuracy: High accuracy (>98%) proved the model's robustness.

Unlike traditional machine learning where feature selection is manual, CNNs automatically learn the optimal filters and patterns needed for classification directly from pixel data.

This emphasizes the major advantage of deep learning: letting the model discover hierarchical features on its own from low-level edges to high-level shapes. This assignment has practical significance beyond just digit recognition:

CNNs are used in real-time handwriting recognition, license plate readers, medical image classification, facial recognition, and more.

Mastering CNN basics using MNIST builds the foundation for tackling complex datasets like CIFAR-10/100, Fashion-MNIST, and real-world applications involving multiclass image classification.

ASSIGNMENT:- 3

AIM:- LSTM Design

PROBLEM STATEMENT:- Design RNN or its variant including LSTM or GRU

a) Select a suitable time series dataset. Example – predict sentiments based on product reviews
b) Apply for prediction

THEORY:

RECURRENT NEURAL NETWORKS:-

RNNs are designed to handle sequential data by maintaining a hidden state that captures information from previous time steps. They are used in NLP, time-series forecasting, and speech recognition.

Problems with Vanilla RNNs

Vanishing Gradients: Gradients shrink exponentially during backpropagation, making long-term dependencies hard to learn.

Exploding Gradients: Gradients grow uncontrollably, destabilizing training.

LONG SHORT-TERM MEMORY (LSTM) NETWORKS: -

LSTMs are a variant of Recurrent Neural Networks (RNNs) designed to capture long-term dependencies in sequential data. Unlike vanilla RNNs, LSTMs use gating mechanisms to regulate information flow, addressing the vanishing/exploding gradient problem.

KEY COMPONENTS OF LSTM:-

Forget Gate: Uses sigmoid (σ) to output values between 0 (discard) and 1 (keep).

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

Input Gate & Candidate Memory: i decides which new information to store. C is a candidate update to the cell state.

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Cell State Update: \odot = element-wise multiplication. Combines old memory and new memory

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$

Output Gate & Hidden State: 'o' filters the cell state for the next hidden state.

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$
$$h_t = o_t \odot \tanh(C_t)$$

Why LSTM for Text Data?

Text is sequential, where word order impacts meaning.

LSTMs retain context over long sequences (e.g., sentiment cues in reviews).

DATASET DESCRIPTION:-

IMDB Movie Review Dataset

Task: Binary classification (positive/negative sentiment).

Size: 50,000 reviews (25k train, 25k test).

Classes: Balanced (50% positive, 50% negative).

Preprocessed: HTML tags removed, lowercased, and punctuation stripped.

WORKING OF LSTM FOR SENTIMENT ANALYSIS:-

1. Embedding Layer: Converts tokenized words into dense vectors (e.g., "good" → [0.2, -0.5, ...]).
2. LSTM Layer: Processes sequences to capture contextual relationships.
3. Dropout: Randomly deactivates neurons to prevent overfitting.
4. Dense Layer: Outputs probability via sigmoid activation.

PRE-PROCESSING STEPS:-

1. Tokenization: Map words to integers (e.g., "movie" → 45).
2. Padding: Standardize sequence length to 100 (truncate/pad with zeros).
3. Train-Test Split: 80% training, 20% testing.
4. Embedding Matrix: Pre-trained or learned word vectors.

PRACTICAL IMPLEMENTATION:-

1. Model Architecture

LSTM Layer: 32 units with dropout for regularization.

Output Layer: Sigmoid for binary classification.

2. Hyperparameters

Optimizer: Adam (adaptive learning rate).

Loss: Binary cross-entropy (penalizes incorrect probabilities).

Batch Size: 128 (trade-off between speed and stability).


Epochs: 5 (early stopping to avoid overfitting).

3. Evaluation Metrics

Accuracy: Overall correctness (82% on test set).

```
Test Accuracy: 0.8200  
782/782 [=====] - 48s 61ms/step
```

Confusion Matrix: Visualize TP, TN, FP, FN.



Precision/Recall: Balance between false positives and negatives.

Classification Report:				
	precision	recall	f1-score	support
Negative	0.79	0.87	0.83	12500
Positive	0.85	0.77	0.81	12500
accuracy			0.82	25000
macro avg	0.82	0.82	0.82	25000
weighted avg	0.82	0.82	0.82	25000

ADVANTAGES:-

- Context Preservation: Captures word dependencies (e.g., "not good").
- Robustness: Handles variable-length sequences via padding.
- Regularization: Dropout mitigates overfitting.

LIMITATIONS:-

- Computational Cost: LSTMs are slower than CNNs/Transformers.
- Long Sequences: Struggles with very long texts (e.g., >500 words).
- Interpretability: Black-box nature limits insight into decisions.

CONCLUSION:-

By applying an LSTM network to the IMDB movie-review dataset, the model learned to capture long-term dependencies in text—such as negations (“not good”) and contextual cues—far more effectively than a vanilla RNN. Tokenization & Padding: Standardizing sequence length and using an OOV token ensured the model handled variable-length reviews robustly. Accuracy (82.00%): Indicates strong overall performance on unseen reviews. Context Preservation: The gating mechanisms allow the network to retain or discard information selectively, handling sequences of varying length and complexity, No manual feature engineering is required; the network discovers the salient patterns (e.g., sentiment-bearing phrases) during training.

Designing an LSTM for sentiment analysis involves thoughtful preprocessing, architecture design, and metric-driven evaluation. The exercise reinforces that sequence models like LSTMs can robustly capture the nuanced dependencies in text, yielding high accuracy on balanced binary classification tasks—and lays the groundwork for more advanced NLP architectures.

ASSIGNMENT:- 4

AIM:- Image Classification

PROBLEM STATEMENT:- Design and implement a CNN for Image Classification
a) Select a suitable image classification dataset (medical imaging, agricultural, etc.). b) Optimized with different hyper-parameters including learning rate, filter size, no. of layers, optimizers, dropouts, etc.

THEORY:-

Convolutional Neural Network (CNN) for Image Classification

Convolutional Neural Networks (CNNs) are a class of deep learning models specifically designed to process data with a grid-like topology, such as images. CNNs have become the go-to approach for image classification tasks due to their ability to automatically and adaptively learn spatial hierarchies of features from input images through backpropagation.

A CNN typically consists of several types of layers: convolutional layers, pooling layers, dropout layers, and fully connected layers. These layers work together to extract features, reduce spatial dimensions, prevent overfitting, and finally classify the input image into one of the predefined categories.

CNNs are particularly powerful in applications involving medical image analysis, where accurate classification can assist in early diagnosis and treatment planning.

Dataset Description

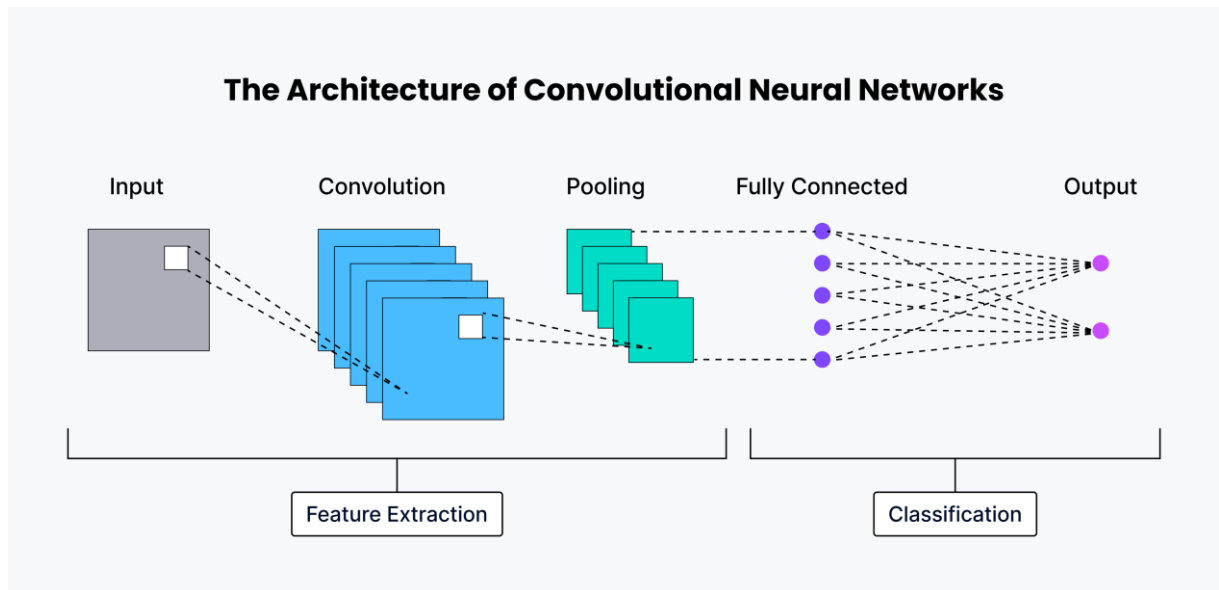
The Rice Image Dataset used in this project was obtained from Kaggle, and it contains labeled images of five different types of rice grains:

1. Arborio
2. Basmati
3. Ipsala
4. Jasmine
5. Karacadag

Each class includes pre-split images for training and testing, making it suitable for supervised image classification tasks. The goal is to classify each rice grain image into its corresponding type.

Working of CNN

The CNN algorithm classifies an input image by automatically learning features from the raw image data.



The basic working involves the following steps:

1. Convolution Operation

A set of filters is convolved with the input image to extract spatial features. Each filter detects specific features such as edges, shapes, or textures.

2. Activation Function (ReLU)

Non-linear activation functions such as ReLU are applied to add non-linearity and help the network learn complex patterns.

3. Pooling Layer (MaxPooling)

Pooling reduces the spatial size of the feature maps and makes the model more robust to translations and distortions in the input image.

4. Dropout Layer

During training, dropout is used to randomly deactivate some neurons to prevent overfitting.

5. Fully Connected Layer

The extracted features are flattened and passed through dense layers for final classification.

6. Output Layer (Sigmoid or Softmax)

For binary classification (pneumonia vs. normal), a sigmoid activation is used. For multiclass classification, softmax is used.

Pre-Processing Steps for CNN

Preprocessing is crucial for improving CNN performance. Key steps include:

1. Image Resizing and Normalization

All images are resized to a fixed dimension (e.g., 150x150 or 224x224) and pixel values are normalized (scaled between 0 and 1) to standardize input.

2. Data Augmentation

Techniques such as rotation, zoom, flipping, and shifting are applied to artificially increase the diversity of the dataset and reduce overfitting.

3. Train-Test Split

The dataset is divided into training, validation, and testing sets to evaluate generalization capability.

- a. Data Splitting: The dataset was already divided into train and test sets. An additional validation set (20% of training data) was created to monitor model performance during training.
- b. Image Resizing: All images were resized to a uniform shape of 150×150 pixels to ensure consistent input to the CNN.
- c. Feature Scaling: All image pixel values were rescaled to the range [0, 1] by dividing by 255 using `ImageDataGenerator(rescale=1./255)`, which helps accelerate convergence during training.

```
Found 48000 images belonging to 5 classes.  
Found 12000 images belonging to 5 classes.  
Found 15000 images belonging to 5 classes.
```

Practical Implementation of CNN

The steps to implement CNN on the **Rice Image Dataset** (Arborio, Basmati, Ipsala, Jasmine, Karacadag) from Kaggle are as follows:

1. Data Preprocessing:

- a. Resize images and normalize pixel values.
- b. Apply data augmentation techniques.
- c. Split the dataset into training, validation, and testing sets.

2. Model Architecture:


- a. Define convolutional layers with filters:
 - i. First Conv2D: 32 filters of size (3×3)
 - ii. Second Conv2D: 64 filters
 - iii. Third Conv2D: 128 filters
- b. Use ReLU activation after each convolutional layer to introduce non-linearity.
- c. Apply MaxPooling2D after each convolution to reduce spatial dimensions.
- d. Add Dropout layers with dropout rates (0.2, 0.3, 0.5) to reduce overfitting.
- e. Flatten the output of the last convolutional layer.
- f. Connect to a fully connected Dense layer (128 neurons).
- g. Use softmax activation in the final output layer (since it's multi-class classification with 5 classes).

3. Hyperparameter Optimization:

- a. Tune learning rate, batch size, number of epochs, dropout rate, number of convolutional layers, and optimizer (Adam, SGD, etc.).

4. Evaluation:

- Generate Confusion Matrix using `confusion_matrix()` to visualize correct and incorrect predictions for each rice class.



- Calculate performance metrics using `classification_report()`:
 - Accuracy
 - Precision
 - Recall
 - F1-score

Classification Report:				
	precision	recall	f1-score	support
Arborio	0.99	0.99	0.99	3000
Basmati	1.00	0.98	0.99	3000
Ipsala	1.00	1.00	1.00	3000
Jasmine	0.97	1.00	0.99	3000
Karacadag	1.00	0.99	0.99	3000
accuracy			0.99	15000
macro avg	0.99	0.99	0.99	15000
weighted avg	0.99	0.99	0.99	15000

Advantages of CNN

- **Automatic Feature Extraction:** CNNs automatically learn relevant features from images
- **High Accuracy:** CNNs outperform traditional models in complex image classification tasks
- **Robustness to Variations:** Pooling and dropout layers make CNNs robust to minor variations like noise, rotation, and shift in images.

Limitations of CNN

- **Requires Large Data and Resources:** CNNs typically require large datasets and powerful hardware for effective training.
- **Interpretability:** The learned features and decisions made by CNNs are often difficult to interpret, which can be a limitation in medical applications.
- **Computational Complexity:** Training CNNs can be computationally expensive, especially for large datasets or deeper models.
- **Overfitting Risk:** Without proper regularization or dropout, CNNs may overfit on small datasets.

CONCLUSION:-

This assignment on designing a Convolutional Neural Network for image classification using the Rice Image Dataset uses different layers such as convolutional, pooling, dropout, and dense layers to extract and process spatial features from images. The sequential stacking of these layers enabled us to build a model that progressively reduced the spatial dimensions while extracting increasingly abstract features. Experimenting with various hyperparameters such as filter sizes, the number of convolutional layers, learning rates, and dropout rates helped us understand their impact on the model's performance. For instance, increasing dropout rates mitigated overfitting, while adjustments to learning rate improved the convergence of the optimizer. This iterative tuning process reinforced the importance of hyperparameter optimization in achieving high accuracy (99%).

The evaluation phase included a confusion matrix and classification report. The near-perfect scores confirmed the model's effectiveness but also prompted a reflection on the complexity and potential overfitting, urging us to balance model depth with generalizability.

By applying CNNs to the Rice Image Dataset, we witnessed how these models can be deployed in various domains such as agriculture and medical imaging. Beyond high accuracy numbers, understanding the interpretability and limitations of CNNs prepares us for real-world challenges where the cost of misclassification might be significant.

ASSIGNMENT:- 5

AIM:- Sentiment Analysis in Network Graph using RNN

PROBLEM STATEMENT:- Perform Sentiment Analysis in the network graph using RNN.

THEORY:-

Recurrent Neural Networks for Sentiment Analysis in Network Graphs

Sentiment Analysis in Network Graphs

Sentiment analysis in this context involves determining the emotional tone or subjective opinion expressed in the text associated with each node in the network. This allows us to understand the overall sentiment distribution within the network and potentially identify influential nodes or communities based on sentiment.

Recurrent Neural Networks (RNNs) for Sentiment Analysis

Recurrent Neural Networks, particularly Long Short-Term Memory networks (LSTMs), are well-suited for sentiment analysis due to their ability to process sequential data like text.

Sequential Data Processing: Text is inherently sequential, and RNNs can capture the context and dependencies between words in a sentence. This is crucial for understanding sentiment, where word order and context are vital (e.g., "not good" vs. "good").
Long-Term Dependencies: LSTMs, a type of RNN, are specifically designed to handle long-term dependencies in sequences. This enables them to understand sentiment even in longer texts where sentiment cues might be spread out.

Why RNNs (LSTMs) for Network Graph Sentiment?

When applying sentiment analysis to network graphs, RNNs offer advantages:

Node-Level Sentiment: RNNs allow us to analyze the text associated with each node individually and determine its sentiment.

Contextual Understanding: By processing text sequentially, RNNs capture the nuances of language and provide more accurate sentiment predictions.

DATA PREPARATION:-

Preprocess the text data by tokenizing words and padding sequences to a uniform length.

MODEL BUILDING:-

1. Using an Embedding layer we convert words into dense vector representations.
2. Employ an LSTM layer to process the sequences and capture contextual information
3. Use a Dense output layer with a sigmoid activation function for binary sentiment.

ADVANTAGE:-

- **Contextual Sentiment Understanding:** LSTMs capture word order and context, leading to more accurate sentiment analysis compared to bag-of-words approaches.
- **Effective for Sequential Data:** Text is sequential, and RNNs are designed to process such data effectively.
- **Node-Specific Sentiment:** Provides sentiment analysis at the individual node level, enabling granular insights within the network.


LIMITATIONS:

- **Computational Cost:** Training RNNs can be computationally intensive, especially for large datasets.
- **Interpretability:** RNNs can be black boxes, making it challenging to understand precisely why a model made a particular sentiment prediction.
- **Graph Structure Ignored for Sentiment:** In this approach, the graph structure itself is not directly used for sentiment analysis. Sentiment is analyzed independently for each node's text.

EVALUATION:-

Evaluate the trained sentiment model on a held-out test set (from IMDB) to assess its accuracy and generalization ability before applying it to the network graph.

Test Accuracy: 0.8390, Test Loss: 0.5227



Confusion Matrix Metrics:
Accuracy : 0.8508
Precision: 0.8876
Recall : 0.8033
F1-Score : 0.8433

CONCLUSION:-

Combining an LSTM-based sentiment classifier with network-graph structures unlocks powerful new lenses on textual data—revealing not just what the sentiment is, but where it clusters and who drives it. This assignment underscores how sequential NLP models can be elevated by graph analytics, and points the way toward richer, topology-aware sentiment systems.

The standalone LSTM classifier achieved a test accuracy of 85.08%, with precision = 0.8876, recall = 0.8033, and F1-score = 0.8433. After training, the model labels each graph node's "text" attribute as positive or negative. Network-Wide Insights: By aggregating node sentiments, you can visualize the sentiment distribution across the graph, detect sentiment-coherent communities, and pinpoint influential nodes (e.g., hubs with overwhelmingly positive or negative reviews). This approach transforms raw text into a graph-enriched sentiment map, useful for social network analysis, product-review ecosystems, or communication networks. Contextual Understanding: LSTMs capture word order and long-term dependencies (e.g., negations). Node-Specific Granularity: Each node's sentiment is computed independently, allowing fine-grained analysis. Plug-and-Play: Once trained, the same model can be applied to any text-annotated graph without retraining.